

Sachin Singh

7041200380 · rajeshsachin786@gmail.com · Prayagraj, Uttar Pradesh · github.com/sachin-11 · portfolio.rasuonline.in

PROFESSIONAL SUMMARY

Senior Full Stack & AI Engineer with **7+ years** delivering production-grade SaaS platforms and Generative AI systems. Built an AI Interview SaaS handling **50–100 sessions/month**, cut LLM API costs by **50%** via intelligent prompt compression and model routing, and shipped **5+ live applications** across e-commerce, healthcare, education, and AI SaaS. Deep expertise in **RAG pipelines, agentic workflows, OpenAI/Groq, Pinecone**, and scalable Node.js/Python backends on AWS.

TECHNICAL SKILLS

GenAI / Agents: OpenAI GPT-4o/mini, Groq LLaMA-3.3, LangChain, LangGraph, RAG, Pinecone, Embeddings, Prompt Engineering, Agentic Workflows

Backend: Node.js, NestJS, FastAPI, Express.js, Python, REST APIs, GraphQL, WebSockets, Cron Jobs

Frontend: Next.js 15/16, React.js, TypeScript, Tailwind CSS

Databases: PostgreSQL, MongoDB, Redis, Prisma ORM, MongoDB Aggregation Pipelines, DB Sharding

Cloud / DevOps: AWS (S3, SNS, RDS, EC2, Amplify, CloudWatch, IAM), Docker, Kubernetes, GitHub Actions, Terraform, CI/CD

Auth / Payments: NextAuth.js, JWT, OAuth (Google), bcrypt, RBAC · Stripe, ACH, Apple Pay, Cash Pay

AI Tools: Playwright, openpyxl, Pillow, BullMQ, SSE, Twilio, Proctoring Systems

PROFESSIONAL EXPERIENCE

Moon Technolabs · Ahmedabad

Senior Full Stack & AI Engineer (GenAI · System Design)

07/2023 – Present

AI Interview Agent — Production GenAI SaaS

- **Contributed backend APIs** for AI-driven interview sessions — candidate evaluation workflows, automated LLM-based feedback generation, and real-time STT/TTS voice interview engine with AI proctoring system.
- **Engineered dual-LLM routing:** Groq LLaMA-3.3-70b (primary) → OpenAI GPT-4o auto-fallback on rate limit with **zero interview interruption**; admin alert on provider switch; **reduced token costs 50%** via prompt compression.
- Built bulk recruitment campaign system: CSV import, candidate comparison dashboard, interview retake flows, S3 audio recording, and automated LLM-based evaluation reports.

Seekrz · Austin Indian · Speech Therapy App — Node.js / NestJS

- **Led full backend architecture for Seekrz** — live marketplace app (App Store + Play Store); designed trade offer engine with 48-hour cron auto-cancel, bulk order APIs, and 4 payment integrations (Stripe, ACH, Apple Pay, Cash Pay).
- Built HIPAA-conscious backend for Speech Therapy Medical App — secure patient records, appointment scheduling, session management using NestJS.
- Developed Austin Indian community platform — restaurant search/filter APIs with geo-based recommendations.

ManekTech · Ahmedabad

Software Engineer — MERN Stack

11/2021 – 05/2023

- **Led backend for Cellcast (Australia)** — high-volume bulk SMS platform; engineered campaign APIs handling thousands of concurrent messages with guaranteed delivery reliability.
- Integrated Stripe subscriptions and ACH direct debit; architected SMS credit wallet system with transactional logging and auto top-up triggers.
- Drove backend optimizations that reduced infrastructure costs and improved API response times under peak load.

Crest Infosystems Pvt. Ltd · Surat

Software Engineer — Node.js

01/2021 – 10/2021

- Built Udemy-style online course platform from scratch — course management, mastermind modules, agency dashboards, and Facebook-based revenue integration.
- Implemented JWT auth, Twilio OTP, Node-Cron automation, and complex MongoDB aggregation pipelines for analytics.

- Designed scalable REST APIs, data retrieval pipelines, and live-class streaming APIs supporting real-time student interactions.

PERSONAL PROJECTS

AI Resume Coach — Next.js 16 · FastAPI · LangGraph · OpenAI GPT-4o · Groq LLaMA-3.3 · Pinecone · PostgreSQL · Stripe · Monaco Editor · TypeScript

resume-ai-coach.rasuonline.in

- **Production GenAI SaaS** — LangGraph multi-agent system with **6 autonomous AI agents**: Resume Improvement (ATS ≥ 70 loop), Candidate Screening (live GitHub API), Learning Path, Panel Interview (3 AI judges), Market Intelligence (salary + demand), and Interview Evaluator (contradiction detection).
- **Real-time voice interview engine** — STT/TTS with persona-based AI interviewers (Google style, Amazon LP, Strict, Friendly); **3-AI Panel Interview** (Technical AI + HR AI + Domain AI simultaneously); Pair Programming Mode with Monaco/VS Code editor and AI code review.
- **Advanced proctoring suite** — tab switch, face detection, noise detection, copy-paste monitoring, real-time confidence analyzer (hesitant/confident/nervous tone); Video Presence Analysis via OpenAI Vision.
- **Recruiter automation suite** — bulk CSV campaigns, ZIP resume upload (50 files), AI Job Match ranking, auto-shortlist pipeline, webhook integrations (Slack, Greenhouse, Lever, Workday); Job Application Agent with gap analysis, cover letter generation (3 tones), and JD auto-tailor.
- **Dual AI provider strategy**: Groq (primary, fast) → OpenAI auto-fallback on rate limit; zero interview interruption; Stripe Free/Pro/Enterprise billing; **50% token cost reduction** via prompt compression.

AI Web Scraper Agent — Playwright · FastAPI · GPT-4o-mini · Next.js · openpyxl · Pillow · Docker

web.scraping.rasuonline.in · github.com/sachin-11/web-scraping-agent

- **Built end-to-end agentic scraper**: async Playwright crawler → Semantic DOM Trimmer (strips nav/footer/scripts) → GPT-4o-mini structured JSON extraction → openpyxl Excel with embedded 80×80px product thumbnails — completes in **under 5 seconds**.
- Engineered **asyncio.gather() concurrent pagination** — 5-page crawls reduced from 20s → 4s (**5× faster**); Semantic Trimming cut LLM context by **50–70%**, slashing API costs.
- Built **MD5-hashed TTL disk cache** (1-hr expiry) — repeated requests return in **<5ms** with zero token spend; two-tier auth gateway prevents unauthorized API access.

AI-Powered HR Policy Assistant — FastAPI · LangGraph · GPT-4o-mini · Pinecone · Next.js 14 · Docker

hr-policy.rasuonline.in

- Implemented **Agentic RAG with auto query-rewriting and document grading** — agent retries retrieval with improved queries when results are irrelevant, significantly boosting answer accuracy over naive RAG.
- Full-stack deployment: Render (FastAPI backend) + Vercel (Next.js frontend) with Docker containerization.

AI Prompt Optimizer — Next.js 15 · OpenAI · Anthropic · Gemini · LangGraph · Redis · BullMQ · PostgreSQL

token.rasuonline.in

- Multi-provider prompt compression tool (OpenAI, Anthropic, Gemini) — LangGraph multi-agent workflow for analysis → compression → review; **SSE streaming** for real-time progress.
- Redis + BullMQ background queue for async processing; PostgreSQL + Prisma for prompt history and cache persistence.

EDUCATION

Bharati Vidyapeeth, Pune — MCA (Master of Computer Applications) · 2017

LANGUAGES

English (Professional) · Hindi (Native)